

CS 205 Project – M4

Parallelizing Crop Phenology
reports via NDVI datasets



Team 18:

Elaine Swanson, Alex Johnson,
Saketh Mynampati, Alyssa Ross

Sequential Baseline:

```
+-----+  
| Functions |  
+-----+
```

```
exec(command)  
downloadFile(query, filename)  
writeDataToFile(data, filename)  
outputProfile(counters, dt)
```

```
+-----+  
| main() |  
+-----+
```

Initialize maps

```
+-----+  
| Loop over years |  
| (2001 to USER_DEFINED) |  
+-----+
```


```
Check for leap years  
Loop over months and days  
Calculate week number  
Iterate over region codes  
Construct download URL and filename  
Download file
```

```
Initialize Geotiff object  
Define functions for reading data  
Initialize profiling variables  
Profile reading files  
Process data and calculate statistics  
Perform SSA
```


```
Write output to files  
Output profiling information
```

End program

One Input Example




Output Example

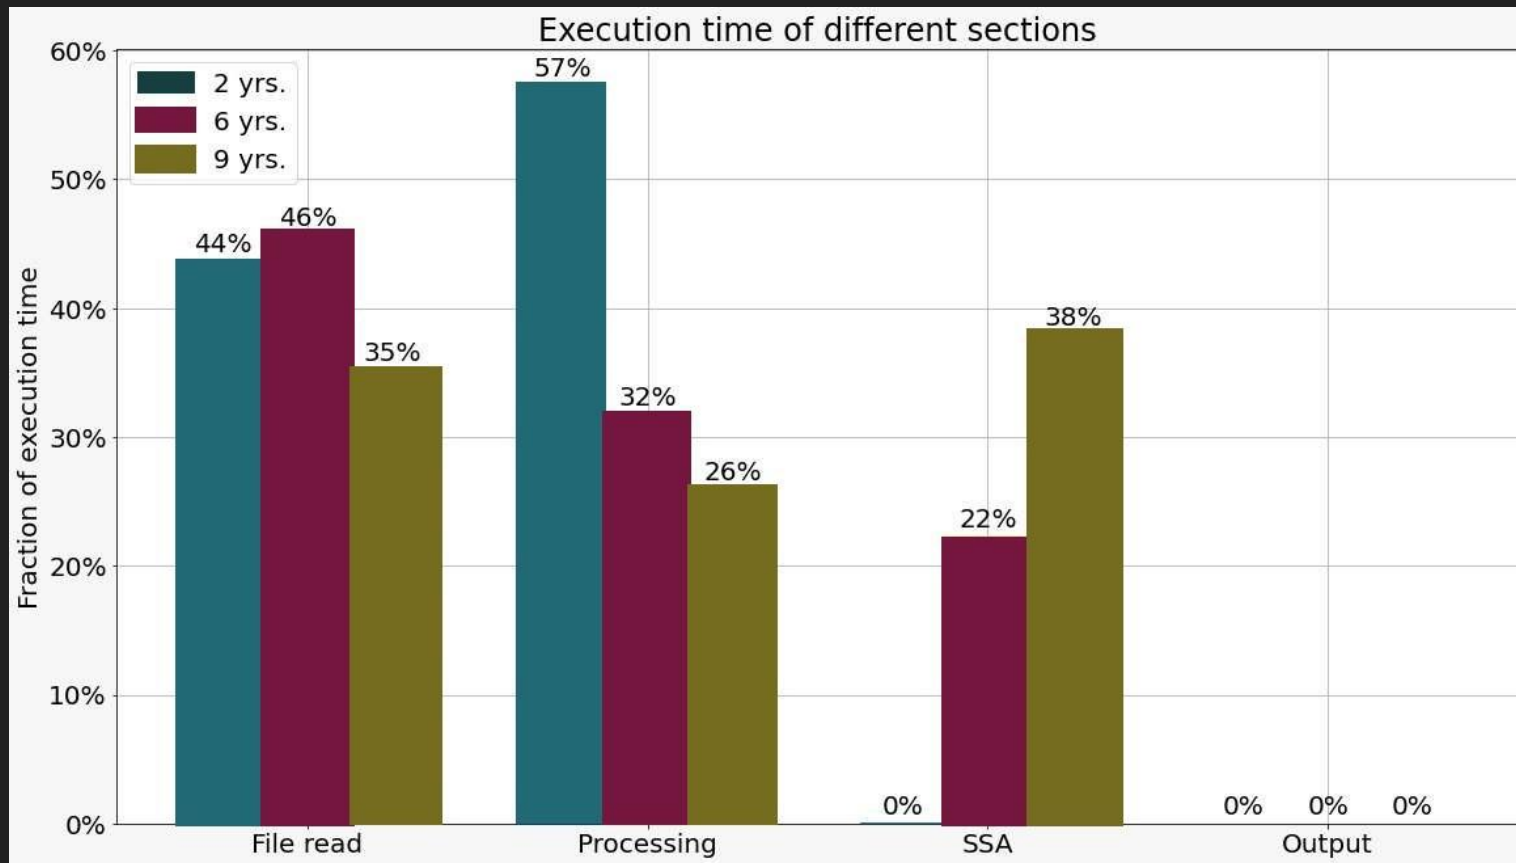


Singular Spectrum Analysis
(SSA) Matrix Binary File
Example: (257, 409)

Output Example:



Sequential Baseline:



==== Processor Information ====

Physical/Logical Cores: 18/36

Frequency (GHz): 3.40

==== Memory Information ====

Total Memory (GB): 68.57

Available Memory (GB): 43.61

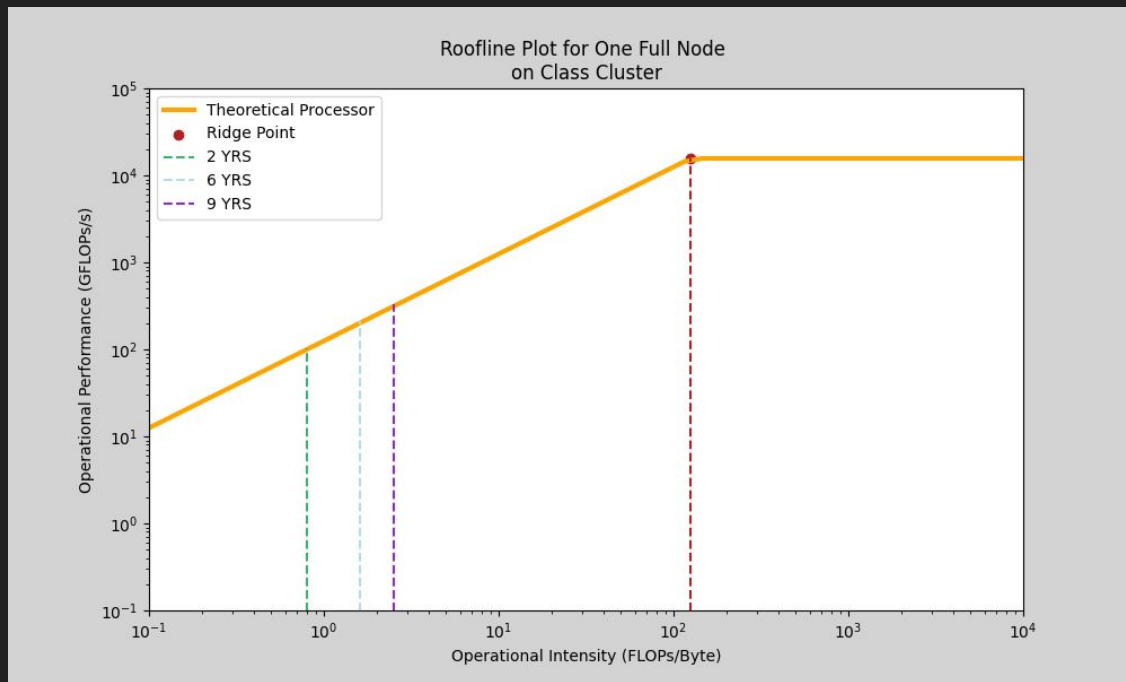
==== Disk Information ====

Device: NVMe storage device

Total Space (GB): 49.99

Free Space (GB): 29.68

salloc -N1 -c36 -t 1:00:00



Peak Floating Point Performance (Integer Precision): 15674 GFLOPs

Peak Memory Bandwidth: 125 GB/s

The ridge point I_{β} is: 125.392

Proposed Parallelization:

- Data: x, y (county) and week/year.
- Threads:
 - Shared disk, shared I/O limitations.
- Main Thread: Coordinates all other OMP threads (36).
 - GOAL: Cropland data is separate from NDVI data, so threads handle loading and overlaying of NDVI and Crop Masks
 - Three kinds of parallelization to test:
 - Simple
 - i. Each thread downloads, overlays, and processes data for a given county, with 1 or 2 threads per county.
 - ii. Two-thread counties communicate via MPI non-blocking operations to minimize waiting - each handles a portion of time.
 - iii. Synchronize at end within county, MPI File Ops for output.
 - Maximize Cache Hit Rate/Resource Use:
 - i. All threads handle single county at a time, and apply mask in sequence.
 - ii. File is resynthesized on single assigned thread via MPI_Reduce for summary statistics over time.
 - iii. Built into time series over chosen years using MPI
 - iv. Heavy MPI communication overhead.
- Note:
 - Program still lies in the memory-bound region of the roofline model.